

rl-project-code-nio

May 4, 2023

```
[1]: #Importing the libraries
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
```

```
[2]: import time
import copy
import numpy as np
import pandas as pd
import chainer
import chainer.functions as F
import chainer.links as L
from plotly import tools
from plotly.graph_objs import *
from plotly.offline import init_notebook_mode, iplot, iplot_mpl
init_notebook_mode()
```

```
[3]: #reading the data
data = pd.read_csv("C:\\Users\\madha\\OneDrive\\Desktop\\RL project\\NIO.csv")
data['Date'] = pd.to_datetime(data['Date'])
data = data.set_index('Date')
print(data.index.min(), data.index.max())
data.head()
```

2020-01-02 00:00:00 2022-12-29 00:00:00

```
[3]:
```

	Open	High	Low	Close	Adj Close	Volume
Date						
2020-01-02	4.10	4.10	3.61	3.72	3.72	103740100
2020-01-03	3.50	3.90	3.48	3.83	3.83	82892400
2020-01-06	4.19	4.24	3.66	3.68	3.68	106619700
2020-01-07	3.70	3.73	3.21	3.24	3.24	106336400
2020-01-08	3.14	3.49	3.13	3.39	3.39	65118100

Splitting the data in train and test model

```
[4]: date_split = '2022-04-01'
train = data[:date_split]
test = data[date_split:]
```

```
len(train), len(test)
```

```
[4]: (568, 188)
```

```
[5]: def plot_train_test(train, test, date_split):

    data = [
        Candlestick(x=train.index, open=train['Open'], high=train['High'],
        ↪ low=train['Low'], close=train['Close'], name='train'),
        Candlestick(x=test.index, open=test['Open'], high=test['High'],
        ↪ low=test['Low'], close=test['Close'], name='test')
    ]
    layout = {
        'shapes': [
            {'x0': date_split, 'x1': date_split, 'y0': 0, 'y1': 1, 'xref':
            ↪ 'x', 'yref': 'paper', 'line': {'color': 'rgb(1,1,0)', 'width': 1}}
        ],
        'annotations': [
            {'x': date_split, 'y': 1.0, 'xref': 'x', 'yref': 'paper',
            ↪ 'showarrow': False, 'xanchor': 'left', 'text': ' test data'},
            {'x': date_split, 'y': 1.0, 'xref': 'x', 'yref': 'paper',
            ↪ 'showarrow': False, 'xanchor': 'right', 'text': 'train data '}
        ]
    }
    figure = Figure(data=data, layout=layout)
    iplot(figure)
    #display(HTML(figure.to_html()))
```

```
[6]: plot_train_test(train, test, date_split)
```

```
[7]: class Environment:

    def __init__(self, data, history_t=90):
        self.data = data
        self.history_t = history_t
        self.reset()

    def reset(self):
        self.t = 0
        self.done = False
        self.profits = 0
        self.positions = []
        self.position_value = 0
        self.history = [0 for _ in range(self.history_t)]
        return [self.position_value] + self.history # obs

    def step(self, act):
```

```

reward = 0

# act = 0: stay, 1: buy, 2: sell
if act == 1:
    self.positions.append(self.data.iloc[self.t, :]['Close'])
elif act == 2: # sell
    if len(self.positions) == 0:
        reward = -1
    else:
        profits = 0
        for p in self.positions:
            profits += (self.data.iloc[self.t, :]['Close'] - p)
        reward += profits
        self.profits += profits
        self.positions = []

# set next time
self.t += 1
self.position_value = 0
for p in self.positions:
    self.position_value += (self.data.iloc[self.t, :]['Close'] - p)
self.history.pop(0)
self.history.append(self.data.iloc[self.t, :]['Close'] - self.data.
↪iloc[(self.t-1), :]['Close'])

# clipping reward
if reward > 0:
    reward = 1
elif reward < 0:
    reward = -1

return [self.position_value] + self.history, reward, self.done # obs, ↪
↪reward, done

```

```

[8]: env = Environment(train)
print(env.reset())
for _ in range(3):
    pact = np.random.randint(3)
    print(env.step(pact))

```

```

[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,

```

[illegible]

Applying the DQN model

```
[9]: # DQN

def train_dqn(env):

    class Q_Network(chainer.Chain):

        def __init__(self, input_size, hidden_size, output_size):
            super(Q_Network, self).__init__(
                L.Linear(input_size, hidden_size),
                L.Linear(hidden_size, hidden_size),
                L.Linear(hidden_size, output_size)
            )

        def __call__(self, x):
            h = F.relu(self.fc1(x))
            h = F.relu(self.fc2(h))
            y = self.fc3(h)
            return y

        def reset(self):
            self.zerograds()

    Q = Q_Network(input_size=env.history_t+1, hidden_size=100, output_size=3)
    Q_ast = copy.deepcopy(Q)
    optimizer = chainer.optimizers.Adam()
    optimizer.setup(Q)

    epoch_num = 50
    step_max = len(env.data)-1
    memory_size = 200
    batch_size = 20
    epsilon = 1.0
    epsilon_decrease = 1e-3
    epsilon_min = 0.1
    start_reduce_epsilon = 200
```

```

train_freq = 10
update_q_freq = 20
gamma = 0.99
show_log_freq = 5

memory = []
total_step = 0
total_rewards = []
total_losses = []

start = time.time()
for epoch in range(epoch_num):

    pobs = env.reset()
    step = 0
    done = False
    total_reward = 0
    total_loss = 0

    while not done and step < step_max:

        # select act
        pact = np.random.randint(3)
        if np.random.rand() > epsilon:
            pact = Q(np.array(pobs, dtype=np.float32).reshape(1, -1))
            pact = np.argmax(pact.data)

        # act
        obs, reward, done = env.step(pact)

        # add memory
        memory.append((pobs, pact, reward, obs, done))
        if len(memory) > memory_size:
            memory.pop(0)

        # train or update q
        if len(memory) == memory_size:
            if total_step % train_freq == 0:
                shuffled_memory = np.random.permutation(memory)
                memory_idx = range(len(shuffled_memory))
                for i in memory_idx[:batch_size]:
                    batch = np.array(shuffled_memory[i:i+batch_size])
                    b_pobs = np.array(batch[:, 0].tolist(), dtype=np.
↪float32).reshape(batch_size, -1)
                    b_pact = np.array(batch[:, 1].tolist(), dtype=np.int32)
                    b_reward = np.array(batch[:, 2].tolist(), dtype=np.
↪int32)

```

```

        b_obs = np.array(batch[:, 3].tolist(), dtype=np.
↪float32).reshape(batch_size, -1)
        b_done = np.array(batch[:, 4].tolist(), dtype=np.bool)

        q = Q(b_pobs)
        maxq = np.max(Q_ast(b_obs).data, axis=1)
        target = copy.deepcopy(q.data)
        for j in range(batch_size):
            target[j, b_pact[j]] =
↪b_reward[j]+gamma*maxq[j]*(not b_done[j])
        Q.reset()
        loss = F.mean_squared_error(q, target)
        total_loss += loss.data
        loss.backward()
        optimizer.update()

        if total_step % update_q_freq == 0:
            Q_ast = copy.deepcopy(Q)

        # epsilon
        if epsilon > epsilon_min and total_step > start_reduce_epsilon:
            epsilon -= epsilon_decrease

        # next step
        total_reward += reward
        pobs = obs
        step += 1
        total_step += 1

    total_rewards.append(total_reward)
    total_losses.append(total_loss)

    if (epoch+1) % show_log_freq == 0:
        log_reward = sum(total_rewards[((epoch+1)-show_log_freq):])/
↪show_log_freq
        log_loss = sum(total_losses[((epoch+1)-show_log_freq):])/
↪show_log_freq
        elapsed_time = time.time()-start
        print('\t'.join(map(str, [epoch+1, epsilon, total_step, log_reward,
↪log_loss, elapsed_time])))
        start = time.time()

    return Q, total_losses, total_rewards

```

```
[10]: Q, total_losses, total_rewards = train_dqn(Environment(train))
```

C:\Users\madha\AppData\Local\Temp\ipykernel_3192\515023404.py:74:

VisibleDeprecationWarning:

Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray.

C:\Users\madha\AppData\Local\Temp\ipykernel_3192\515023404.py:82:

DeprecationWarning:

`np.bool` is a deprecated alias for the builtin `bool`. To silence this warning, use `bool` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.bool_` here. Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

5	0.0999999999999992	2835	-31.0	13367.827018247172
18.211485385894775				
10	0.0999999999999992	5670	-14.6	394245.34012479783
27.150068759918213				
15	0.0999999999999992	8505	-150.4	494851.8175792694
24.27080273628235				
20	0.0999999999999992	11340	-138.6	53307.133459949495
22.52271318435669				
25	0.0999999999999992	14175	-124.4	18333.73999013901
22.55620002746582				
30	0.0999999999999992	17010	-87.4	13681.400502169132
22.050071477890015				
35	0.0999999999999992	19845	-80.6	6073.248950099945
22.67339825630188				
40	0.0999999999999992	22680	-63.4	3513.8264394462108
23.218868255615234				
45	0.0999999999999992	25515	-49.0	2317.257956920564
23.792641401290894				
50	0.0999999999999992	28350	-43.0	1462.8681479632855
23.048948049545288				

```
[11]: def plot_loss_reward(total_losses, total_rewards):  
  
    figure = tools.make_subplots(rows=1, cols=2, subplot_titles=('DQN Loss', 'DQN Reward'), print_grid=False)  
    figure.append_trace(Scatter(y=total_losses, mode='lines', line=dict(color='skyblue')), 1, 1)  
    figure.append_trace(Scatter(y=total_rewards, mode='lines', line=dict(color='orange')), 1, 2)  
    figure['layout']['xaxis1'].update(title='epoch')  
    figure['layout']['xaxis2'].update(title='epoch')  
    figure['layout'].update(height=400, width=900, showlegend=False)
```

```
ipplot(figure)
```

```
[12]: plot_loss_reward(total_losses, total_rewards)
```

C:\Users\madha\anaconda3\lib\site-packages\plotly\tools.py:461:

DeprecationWarning:

plotly.tools.make_subplots is deprecated, please use
plotly.subplots.make_subplots instead

```
[13]: def plot_train_test_by_q(train_env, test_env, Q, algorithm_name):
```

```
    # train
```

```
    pobs = train_env.reset()
```

```
    train_acts = []
```

```
    train_rewards = []
```

```
    for _ in range(len(train_env.data)-1):
```

```
        pact = Q(np.array(pobs, dtype=np.float32).reshape(1, -1))
```

```
        pact = np.argmax(pact.data)
```

```
        train_acts.append(pact)
```

```
        obs, reward, done = train_env.step(pact)
```

```
        train_rewards.append(reward)
```

```
        pobs = obs
```

```
    train_profits = train_env.profits
```

```
    # test
```

```
    pobs = test_env.reset()
```

```
    test_acts = []
```

```
    test_rewards = []
```

```
    for _ in range(len(test_env.data)-1):
```

```
        pact = Q(np.array(pobs, dtype=np.float32).reshape(1, -1))
```

```
        pact = np.argmax(pact.data)
```

```
        test_acts.append(pact)
```

```
        obs, reward, done = test_env.step(pact)
```

```
        test_rewards.append(reward)
```

```
        pobs = obs
```



```

test_profits = test_env.profits

# plot
train_copy = train_env.data.copy()
test_copy = test_env.data.copy()
train_copy['act'] = train_acts + [np.nan]
train_copy['reward'] = train_rewards + [np.nan]
test_copy['act'] = test_acts + [np.nan]
test_copy['reward'] = test_rewards + [np.nan]
train0 = train_copy[train_copy['act'] == 0]
train1 = train_copy[train_copy['act'] == 1]
train2 = train_copy[train_copy['act'] == 2]
test0 = test_copy[test_copy['act'] == 0]
test1 = test_copy[test_copy['act'] == 1]
test2 = test_copy[test_copy['act'] == 2]
act_color0, act_color1, act_color2 = 'black', 'green', 'red'

data = [
    Candlestick(x=train0.index, open=train0['Open'], high=train0['High'],
    ↪low=train0['Low'], close=train0['Close'],
    ↪increasing=dict(line=dict(color=act_color0)),
    ↪decreasing=dict(line=dict(color=act_color0))),
    Candlestick(x=train1.index, open=train1['Open'], high=train1['High'],
    ↪low=train1['Low'], close=train1['Close'],
    ↪increasing=dict(line=dict(color=act_color1)),
    ↪decreasing=dict(line=dict(color=act_color1))),
    Candlestick(x=train2.index, open=train2['Open'], high=train2['High'],
    ↪low=train2['Low'], close=train2['Close'],
    ↪increasing=dict(line=dict(color=act_color2)),
    ↪decreasing=dict(line=dict(color=act_color2))),
    Candlestick(x=test0.index, open=test0['Open'], high=test0['High'],
    ↪low=test0['Low'], close=test0['Close'],
    ↪increasing=dict(line=dict(color=act_color0)),
    ↪decreasing=dict(line=dict(color=act_color0))),
    Candlestick(x=test1.index, open=test1['Open'], high=test1['High'],
    ↪low=test1['Low'], close=test1['Close'],
    ↪increasing=dict(line=dict(color=act_color1)),
    ↪decreasing=dict(line=dict(color=act_color1))),
    Candlestick(x=test2.index, open=test2['Open'], high=test2['High'],
    ↪low=test2['Low'], close=test2['Close'],
    ↪increasing=dict(line=dict(color=act_color2)),
    ↪decreasing=dict(line=dict(color=act_color2)))
]
title = '{}: train s-reward {}, profits {}, test s-reward {}, profits {}'.
↪format(
    algorithm_name,

```

```

        int(sum(train_rewards)),
        int(train_profits),
        int(sum(test_rewards)),
        int(test_profits)
    )
    layout = {
        'title': title,
        'showlegend': False,
        'shapes': [
            {'x0': date_split, 'x1': date_split, 'y0': 0, 'y1': 1, 'xref': 'x', 'yref': 'paper', 'line': {'color': 'rgb(0,0,0)', 'width': 1}},
        ],
        'annotations': [
            {'x': date_split, 'y': 1.0, 'xref': 'x', 'yref': 'paper', 'showarrow': False, 'xanchor': 'left', 'text': ' test data'},
            {'x': date_split, 'y': 1.0, 'xref': 'x', 'yref': 'paper', 'showarrow': False, 'xanchor': 'right', 'text': 'train data '}
        ]
    }
    figure = Figure(data=data, layout=layout)
    iplot(figure)

```

```
[14]: plot_train_test_by_q(Environment(train), Environment(test), Q, 'DQN')
```

Applying the Double DQN model

```

[15]: # Double DQN

def train_ddqn(env):

    class Q_Network(chainer.Chain):

        def __init__(self, input_size, hidden_size, output_size):
            super(Q_Network, self).__init__(
                fc1 = L.Linear(input_size, hidden_size),
                fc2 = L.Linear(hidden_size, hidden_size),
                fc3 = L.Linear(hidden_size, output_size)
            )

        def __call__(self, x):
            h = F.relu(self.fc1(x))
            h = F.relu(self.fc2(h))
            y = self.fc3(h)
            return y

        def reset(self):
            self.zerograds()

```

```

Q = Q_Network(input_size=env.history_t+1, hidden_size=100, output_size=3)
Q_ast = copy.deepcopy(Q)
optimizer = chainer.optimizers.Adam()
optimizer.setup(Q)

epoch_num = 50
step_max = len(env.data)-1
memory_size = 200
batch_size = 50
epsilon = 1.0
epsilon_decrease = 1e-3
epsilon_min = 0.1
start_reduce_epsilon = 200
train_freq = 10
update_q_freq = 20
gamma = 0.99
show_log_freq = 5

memory = []
total_step = 0
total_rewards = []
total_losses = []

start = time.time()
for epoch in range(epoch_num):

    pobs = env.reset()
    step = 0
    done = False
    total_reward = 0
    total_loss = 0

    while not done and step < step_max:

        # select act
        pact = np.random.randint(3)
        if np.random.rand() > epsilon:
            pact = Q(np.array(pobs, dtype=np.float32).reshape(1, -1))
            pact = np.argmax(pact.data)

        # act
        obs, reward, done = env.step(pact)

        # add memory
        memory.append((pobs, pact, reward, obs, done))
        if len(memory) > memory_size:

```

```

memory.pop(0)

# train or update q
if len(memory) == memory_size:
    if total_step % train_freq == 0:
        shuffled_memory = np.random.permutation(memory)
        memory_idx = range(len(shuffled_memory))
        for i in memory_idx[:batch_size]:
            batch = np.array(shuffled_memory[i:i+batch_size])
            b_pobs = np.array(batch[:, 0].tolist(), dtype=np.
↪float32).reshape(batch_size, -1)
            b_pact = np.array(batch[:, 1].tolist(), dtype=np.int32)
            b_reward = np.array(batch[:, 2].tolist(), dtype=np.
↪int32)
            b_obs = np.array(batch[:, 3].tolist(), dtype=np.
↪float32).reshape(batch_size, -1)
            b_done = np.array(batch[:, 4].tolist(), dtype=np.bool)

            q = Q(b_pobs)
            """ <<< DQN -> Double DQN
            maxq = np.max(Q_ast(b_obs).data, axis=1)
            == """
            indices = np.argmax(q.data, axis=1)
            maxqs = Q_ast(b_obs).data
            """ >>> """
            target = copy.deepcopy(q.data)
            for j in range(batch_size):
                """ <<< DQN -> Double DQN
                target[j, b_pact[j]] =
↪b_reward[j]+gamma*maxq[j]*(not b_done[j])
                == """
                target[j, b_pact[j]] = b_reward[j]+gamma*maxqs[j,
↪indices[j]]*(not b_done[j])
                """ >>> """
            Q.reset()
            loss = F.mean_squared_error(q, target)
            total_loss += loss.data
            loss.backward()
            optimizer.update()

        if total_step % update_q_freq == 0:
            Q_ast = copy.deepcopy(Q)

# epsilon
if epsilon > epsilon_min and total_step > start_reduce_epsilon:
    epsilon -= epsilon_decrease

```

```

        # next step
        total_reward += reward
        pobs = obs
        step += 1
        total_step += 1

    total_rewards.append(total_reward)
    total_losses.append(total_loss)

    if (epoch+1) % show_log_freq == 0:
        log_reward = sum(total_rewards[((epoch+1)-show_log_freq):])/
↪show_log_freq
        log_loss = sum(total_losses[((epoch+1)-show_log_freq):])/
↪show_log_freq
        elapsed_time = time.time()-start
        print('\t'.join(map(str, [epoch+1, epsilon, total_step, log_reward,
↪log_loss, elapsed_time])))
        start = time.time()

    return Q, total_losses, total_rewards

```

```
[16]: Q, total_losses, total_rewards = train_ddqn(Environment(train))
```

C:\Users\madha\AppData\Local\Temp\ipykernel_3192\1579113142.py:74:
VisibleDeprecationWarning:

Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray.

C:\Users\madha\AppData\Local\Temp\ipykernel_3192\1579113142.py:82:
DeprecationWarning:

`np.bool` is a deprecated alias for the builtin `bool`. To silence this warning, use `bool` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.bool_` here.
Deprecated in NumPy 1.20; for more details and guidance:
<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

5	0.09999999999999992	2835	-33.4	41.554668059386316
14.612796068191528				
10	0.09999999999999992	5670	3.6	79.02428655717522
14.437743902206421				
15	0.09999999999999992	8505	4.2	93.92526926249266
15.615903377532959				
20	0.09999999999999992	11340	4.6	14.700800946727394
13.950337886810303				

25	0.09999999999999992	14175	15.4	7.628558929357678
30	0.09999999999999992	17010	13.8	6.4270844518207015
35	0.09999999999999992	19845	10.8	8.912814674107358
40	0.09999999999999992	22680	15.8	7.5310753562022
45	0.09999999999999992	25515	16.4	6.759689093055203
50	0.09999999999999992	28350	11.4	5.152364396560006

```
[17]: def plot_loss_reward(total_losses, total_rewards):

    figure = tools.make_subplots(rows=1, cols=2, subplot_titles=(' Double DQN Loss', ' Double DQN Reward'), print_grid=False)
    figure.append_trace(Scatter(y=total_losses, mode='lines', line=dict(color='skyblue')), 1, 1)
    figure.append_trace(Scatter(y=total_rewards, mode='lines', line=dict(color='orange')), 1, 2)
    figure['layout']['xaxis1'].update(title='epoch')
    figure['layout']['xaxis2'].update(title='epoch')
    figure['layout'].update(height=400, width=900, showlegend=False)
    iplot(figure)
```

```
[18]: plot_loss_reward(total_losses, total_rewards)
```

C:\Users\madha\anaconda3\lib\site-packages\plotly\tools.py:461:
DeprecationWarning:

plotly.tools.make_subplots is deprecated, please use
plotly.subplots.make_subplots instead

```
[19]: plot_train_test_by_q(Environment(train), Environment(test), Q, 'Double DQN')
```

```
[ ]:
```